

# ECE 285 Assignment 1: Classification using Neural Network

Now that you have developed and tested your model on the toy dataset set. It's time to get down and get dirty with a standard dataset such as cifar10. At this point, you will be using the provided training data to tune the hyper-parameters of your network such that it works with cifar10 for the task of multi-class classification.

Important: Recall that now we have non-linear decision boundaries, thus we do not need to do one vs all classification. We learn a single non-linear decision boundary instead. Our non-linear boundaries (thanks to relu non-linearity) will take care of differentiating between all the classes

TO SUBMIT: PDF of this notebook with all the required outputs and answers.

```
In [2]: # Prepare Packages
import numpy as np
import matplotlib.pyplot as plt

from ece285.utils.data_processing import get_cifar10_data
from ece285.utils.evaluation import get_classification_accuracy

%matplotlib inline
plt.rcParams["figure.figsize"] = (10.0, 8.0) # set default size of plots

# For auto-reloading external modules
# See http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2

# Use a subset of CIFAR10 for the assignment
dataset = get_cifar10_data(
    subset_train=5000,
    subset_val=250,
    subset_test=500,
)

print(dataset.keys())
print("Training Set Data Shape: ", dataset["x_train"].shape)
print("Training Set Label Shape: ", dataset["y_train"].shape)
print("Validation Set Data Shape: ", dataset["x_val"].shape)
print("Validation Set Label Shape: ", dataset["y_val"].shape)
print("Test Set Data Shape: ", dataset["x_test"].shape)
print("Test Set Label Shape: ", dataset["y_test"].shape)
```

```
The autoreload extension is already loaded. To reload it, use:  
    %reload_ext autoreload  
dict_keys(['x_train', 'y_train', 'x_val', 'y_val', 'x_test', 'y_test'])  
Training Set Data Shape: (5000, 3072)  
Training Set Label Shape: (5000,)  
Validation Set Data Shape: (250, 3072)  
Validation Set Label Shape: (250,)  
Test Set Data Shape: (500, 3072)  
Test Set Label Shape: (500,)
```

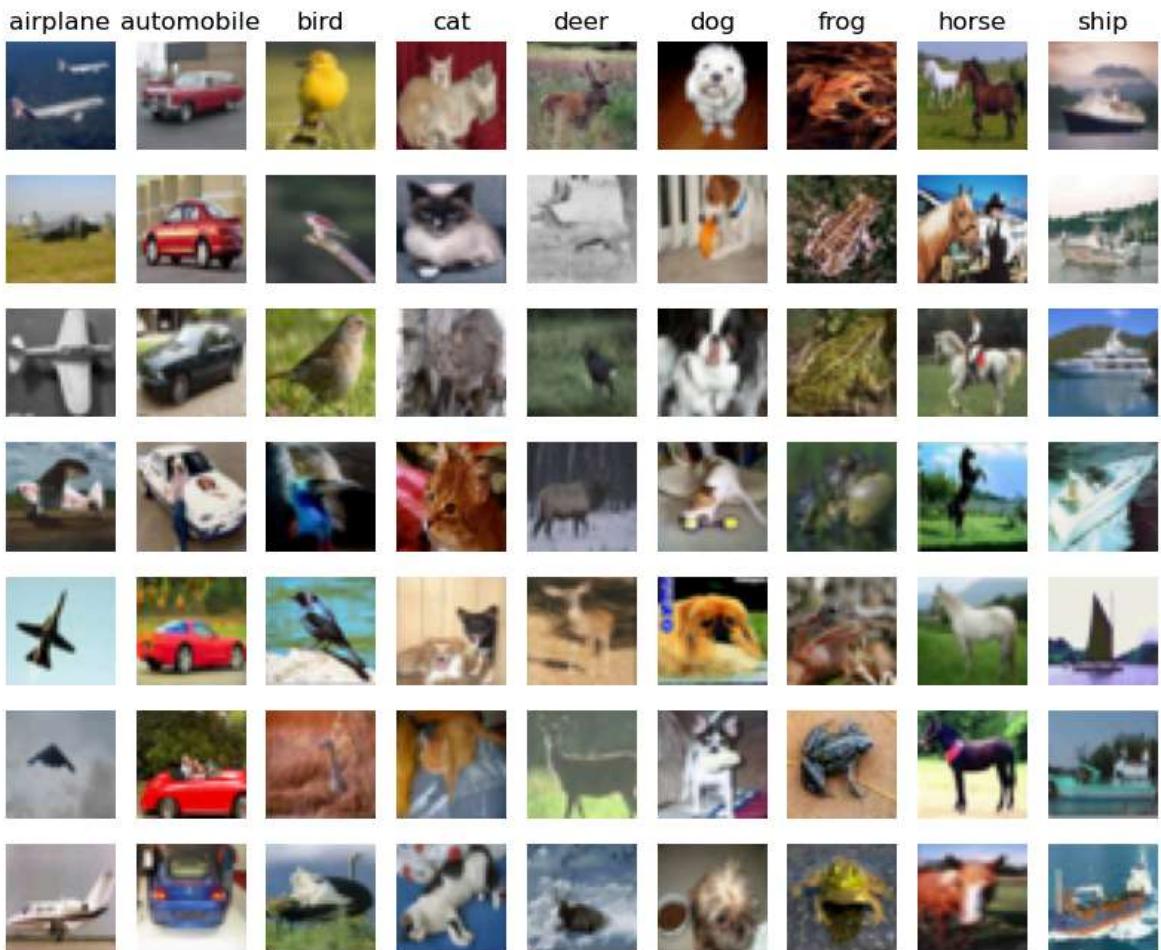
```
In [4]: x_train = dataset["x_train"]  
y_train = dataset["y_train"]  
x_val = dataset["x_val"]  
y_val = dataset["y_val"]  
x_test = dataset["x_test"]  
y_test = dataset["y_test"]
```

```
In [5]: # Import more utilities and the Layers you have implemented  
from ece285.layers.sequential import Sequential  
from ece285.layers.linear import Linear  
from ece285.layers.relu import ReLU  
from ece285.layers.softmax import Softmax  
from ece285.layers.loss_func import CrossEntropyLoss  
from ece285.utils.optimizer import SGD  
from ece285.utils.dataset import DataLoader  
from ece285.utils.trainer import Trainer
```

## Visualize some examples from the dataset.

```
In [6]: # We show a few examples of training images from each class.  
classes = [  
    "airplane",  
    "automobile",  
    "bird",  
    "cat",  
    "deer",  
    "dog",  
    "frog",  
    "horse",  
    "ship",  
]  
samples_per_class = 7  
  
  
def visualize_data(dataset, classes, samples_per_class):  
    num_classes = len(classes)  
    for y, cls in enumerate(classes):  
        idxs = np.flatnonzero(y_train == y)  
        idxs = np.random.choice(idxs, samples_per_class, replace=False)  
        for i, idx in enumerate(idxs):  
            plt_idx = i * num_classes + y + 1  
            plt.subplot(samples_per_class, num_classes, plt_idx)  
            plt.imshow(dataset[idx])  
            plt.axis("off")  
            if i == 0:  
                plt.title(cls)  
    plt.show()
```

```
# Visualize the first 10 classes
visualize_data(
    x_train.reshape(5000, 3, 32, 32).transpose(0, 2, 3, 1),
    classes,
    samples_per_class,
)
```



## Initialize the model

```
In [7]: input_size = 3072
hidden_size = 100 # Hidden Layer size (Hyper-parameter)
num_classes = 10 # Output

# For a default setting we use the same model we used for the toy dataset.
# This tells you the power of a 2 Layered Neural Network. Recall the Universal A
# A 2 Layer neural network with non-linearities can approximate any function, gi
def init_model():
    # np.random.seed(0) # No need to fix the seed here
    l1 = Linear(input_size, hidden_size)
    l2 = Linear(hidden_size, num_classes)

    r1 = ReLU()
    softmax = Softmax()
    return Sequential([l1, r1, l2, softmax])
```

```
In [8]: # Initialize the dataset with the dataloader class
dataset = DataLoader(x_train, y_train, x_val, y_val, x_test, y_test)
net = init_model()
```

```
optim = SGD(net, lr=0.01, weight_decay=0.01)
loss_func = CrossEntropyLoss()
epoch = 200 # (Hyper-parameter)
batch_size = 200 # (Reduce the batch size if your computer is unable to handle
```

```
In [9]: # Initialize the trainer class by passing the above modules
trainer = Trainer(
    dataset, optim, net, loss_func, epoch, batch_size, validate_interval=3
)
```

```
In [10]: # Call the trainer function we have already implemented for you. This trains the
# hyper-parameters. It follows the same procedure as in the last ipython notebook
train_error, validation_accuracy = trainer.train()
```

Epoch Average Loss: 2.302537  
Validate Acc: 0.084  
Epoch Average Loss: 2.302357  
Epoch Average Loss: 2.302150  
Epoch Average Loss: 2.301861  
Validate Acc: 0.108  
Epoch Average Loss: 2.301441  
Epoch Average Loss: 2.300843  
Epoch Average Loss: 2.299991  
Validate Acc: 0.100  
Epoch Average Loss: 2.298838  
Epoch Average Loss: 2.297339  
Epoch Average Loss: 2.295524  
Validate Acc: 0.084  
Epoch Average Loss: 2.293417  
Epoch Average Loss: 2.290900  
Epoch Average Loss: 2.287854  
Validate Acc: 0.084  
Epoch Average Loss: 2.283963  
Epoch Average Loss: 2.278960  
Epoch Average Loss: 2.272802  
Validate Acc: 0.096  
Epoch Average Loss: 2.265813  
Epoch Average Loss: 2.258376  
Epoch Average Loss: 2.250734  
Validate Acc: 0.104  
Epoch Average Loss: 2.243073  
Epoch Average Loss: 2.235618  
Epoch Average Loss: 2.228595  
Validate Acc: 0.120  
Epoch Average Loss: 2.221881  
Epoch Average Loss: 2.215906  
Epoch Average Loss: 2.210026  
Validate Acc: 0.124  
Epoch Average Loss: 2.204771  
Epoch Average Loss: 2.199859  
Epoch Average Loss: 2.195318  
Validate Acc: 0.136  
Epoch Average Loss: 2.191022  
Epoch Average Loss: 2.187192  
Epoch Average Loss: 2.183369  
Validate Acc: 0.140  
Epoch Average Loss: 2.179967  
Epoch Average Loss: 2.176628  
Epoch Average Loss: 2.173439  
Validate Acc: 0.144  
Epoch Average Loss: 2.169992  
Epoch Average Loss: 2.167227  
Epoch Average Loss: 2.164560  
Validate Acc: 0.140  
Epoch Average Loss: 2.161789  
Epoch Average Loss: 2.159188  
Epoch Average Loss: 2.156697  
Validate Acc: 0.144  
Epoch Average Loss: 2.154358  
Epoch Average Loss: 2.152055  
Epoch Average Loss: 2.150305  
Validate Acc: 0.148  
Epoch Average Loss: 2.147913  
Epoch Average Loss: 2.145961

Epoch Average Loss: 2.144044  
Validate Acc: 0.148  
Epoch Average Loss: 2.142169  
Epoch Average Loss: 2.140343  
Epoch Average Loss: 2.138655  
Validate Acc: 0.152  
Epoch Average Loss: 2.136682  
Epoch Average Loss: 2.135321  
Epoch Average Loss: 2.133533  
Validate Acc: 0.152  
Epoch Average Loss: 2.132009  
Epoch Average Loss: 2.130416  
Epoch Average Loss: 2.129122  
Validate Acc: 0.152  
Epoch Average Loss: 2.127431  
Epoch Average Loss: 2.125651  
Epoch Average Loss: 2.124258  
Validate Acc: 0.168  
Epoch Average Loss: 2.122828  
Epoch Average Loss: 2.121352  
Epoch Average Loss: 2.119664  
Validate Acc: 0.156  
Epoch Average Loss: 2.118078  
Epoch Average Loss: 2.116450  
Epoch Average Loss: 2.114914  
Validate Acc: 0.172  
Epoch Average Loss: 2.112643  
Epoch Average Loss: 2.110706  
Epoch Average Loss: 2.109698  
Validate Acc: 0.180  
Epoch Average Loss: 2.107314  
Epoch Average Loss: 2.104992  
Epoch Average Loss: 2.103629  
Validate Acc: 0.172  
Epoch Average Loss: 2.101204  
Epoch Average Loss: 2.098445  
Epoch Average Loss: 2.096262  
Validate Acc: 0.180  
Epoch Average Loss: 2.093518  
Epoch Average Loss: 2.091024  
Epoch Average Loss: 2.088212  
Validate Acc: 0.184  
Epoch Average Loss: 2.085072  
Epoch Average Loss: 2.082427  
Epoch Average Loss: 2.079305  
Validate Acc: 0.228  
Epoch Average Loss: 2.076111  
Epoch Average Loss: 2.073149  
Epoch Average Loss: 2.069611  
Validate Acc: 0.224  
Epoch Average Loss: 2.066161  
Epoch Average Loss: 2.063133  
Epoch Average Loss: 2.060298  
Validate Acc: 0.228  
Epoch Average Loss: 2.056917  
Epoch Average Loss: 2.053934  
Epoch Average Loss: 2.050681  
Validate Acc: 0.228  
Epoch Average Loss: 2.047618  
Epoch Average Loss: 2.044920

Epoch Average Loss: 2.041692  
Validate Acc: 0.248  
Epoch Average Loss: 2.038495  
Epoch Average Loss: 2.036630  
Epoch Average Loss: 2.033501  
Validate Acc: 0.248  
Epoch Average Loss: 2.031365  
Epoch Average Loss: 2.028272  
Epoch Average Loss: 2.025925  
Validate Acc: 0.264  
Epoch Average Loss: 2.023761  
Epoch Average Loss: 2.021007  
Epoch Average Loss: 2.018870  
Validate Acc: 0.244  
Epoch Average Loss: 2.016123  
Epoch Average Loss: 2.014135  
Epoch Average Loss: 2.011489  
Validate Acc: 0.268  
Epoch Average Loss: 2.009680  
Epoch Average Loss: 2.008137  
Epoch Average Loss: 2.005660  
Validate Acc: 0.264  
Epoch Average Loss: 2.003601  
Epoch Average Loss: 2.001001  
Epoch Average Loss: 1.999229  
Validate Acc: 0.264  
Epoch Average Loss: 1.997444  
Epoch Average Loss: 1.995157  
Epoch Average Loss: 1.993021  
Validate Acc: 0.272  
Epoch Average Loss: 1.991517  
Epoch Average Loss: 1.989319  
Epoch Average Loss: 1.987260  
Validate Acc: 0.276  
Epoch Average Loss: 1.985221  
Epoch Average Loss: 1.983659  
Epoch Average Loss: 1.981298  
Validate Acc: 0.284  
Epoch Average Loss: 1.978509  
Epoch Average Loss: 1.976127  
Epoch Average Loss: 1.973707  
Validate Acc: 0.292  
Epoch Average Loss: 1.972199  
Epoch Average Loss: 1.969201  
Epoch Average Loss: 1.966288  
Validate Acc: 0.288  
Epoch Average Loss: 1.963509  
Epoch Average Loss: 1.961164  
Epoch Average Loss: 1.958564  
Validate Acc: 0.296  
Epoch Average Loss: 1.955302  
Epoch Average Loss: 1.951548  
Epoch Average Loss: 1.948651  
Validate Acc: 0.300  
Epoch Average Loss: 1.945868  
Epoch Average Loss: 1.942914  
Epoch Average Loss: 1.939747  
Validate Acc: 0.268  
Epoch Average Loss: 1.936792  
Epoch Average Loss: 1.934389

Epoch Average Loss: 1.931108  
Validate Acc: 0.292  
Epoch Average Loss: 1.927453  
Epoch Average Loss: 1.926925  
Epoch Average Loss: 1.924990  
Validate Acc: 0.284  
Epoch Average Loss: 1.921991  
Epoch Average Loss: 1.919087  
Epoch Average Loss: 1.917199  
Validate Acc: 0.284  
Epoch Average Loss: 1.914289  
Epoch Average Loss: 1.911488  
Epoch Average Loss: 1.910590  
Validate Acc: 0.288  
Epoch Average Loss: 1.907297  
Epoch Average Loss: 1.906047  
Epoch Average Loss: 1.902912  
Validate Acc: 0.304  
Epoch Average Loss: 1.901049  
Epoch Average Loss: 1.899232  
Epoch Average Loss: 1.896585  
Validate Acc: 0.296  
Epoch Average Loss: 1.894317  
Epoch Average Loss: 1.892472  
Epoch Average Loss: 1.889595  
Validate Acc: 0.316  
Epoch Average Loss: 1.887675  
Epoch Average Loss: 1.885574  
Epoch Average Loss: 1.884415  
Validate Acc: 0.300  
Epoch Average Loss: 1.882466  
Epoch Average Loss: 1.879865  
Epoch Average Loss: 1.877374  
Validate Acc: 0.288  
Epoch Average Loss: 1.876498  
Epoch Average Loss: 1.872763  
Epoch Average Loss: 1.872313  
Validate Acc: 0.312  
Epoch Average Loss: 1.868820  
Epoch Average Loss: 1.867578  
Epoch Average Loss: 1.865027  
Validate Acc: 0.320  
Epoch Average Loss: 1.863167  
Epoch Average Loss: 1.861850  
Epoch Average Loss: 1.859246  
Validate Acc: 0.316  
Epoch Average Loss: 1.857995  
Epoch Average Loss: 1.855283  
Epoch Average Loss: 1.853868  
Validate Acc: 0.316  
Epoch Average Loss: 1.850525  
Epoch Average Loss: 1.850173  
Epoch Average Loss: 1.847342  
Validate Acc: 0.320  
Epoch Average Loss: 1.844937  
Epoch Average Loss: 1.842297  
Epoch Average Loss: 1.841978  
Validate Acc: 0.300  
Epoch Average Loss: 1.840899  
Epoch Average Loss: 1.837400

```
Epoch Average Loss: 1.835204
Validate Acc: 0.308
Epoch Average Loss: 1.836128
Epoch Average Loss: 1.832263
Epoch Average Loss: 1.831628
Validate Acc: 0.312
Epoch Average Loss: 1.828670
Epoch Average Loss: 1.825933
Epoch Average Loss: 1.824297
Validate Acc: 0.336
Epoch Average Loss: 1.822668
Epoch Average Loss: 1.822196
Epoch Average Loss: 1.820392
Validate Acc: 0.312
Epoch Average Loss: 1.815700
Epoch Average Loss: 1.816024
Epoch Average Loss: 1.811975
Validate Acc: 0.328
Epoch Average Loss: 1.811493
Epoch Average Loss: 1.808913
Epoch Average Loss: 1.807433
Validate Acc: 0.324
Epoch Average Loss: 1.807012
Epoch Average Loss: 1.804324
Epoch Average Loss: 1.803040
Validate Acc: 0.336
Epoch Average Loss: 1.801166
```

Print the training and validation accuracies for the default hyper-parameters provided

```
In [11]: from ece285.utils.evaluation import get_classification_accuracy
```

```
out_train = net.predict(x_train)
acc = get_classification_accuracy(out_train, y_train)
print("Training acc: ", acc)
out_val = net.predict(x_val)
acc = get_classification_accuracy(out_val, y_val)
print("Validation acc: ", acc)
```

```
Training acc: 0.35
Validation acc: 0.316
```

## Debug the training

With the default parameters we provided above, you should get a validation accuracy of around ~0.2 on the validation set. This isn't very good.

One strategy for getting insight into what's wrong is to plot the training loss function and the validation accuracies during optimization.

Another strategy is to visualize the weights that were learned in the first layer of the network. In most neural networks trained on visual data, the first layer weights typically show some visible structure when visualized.

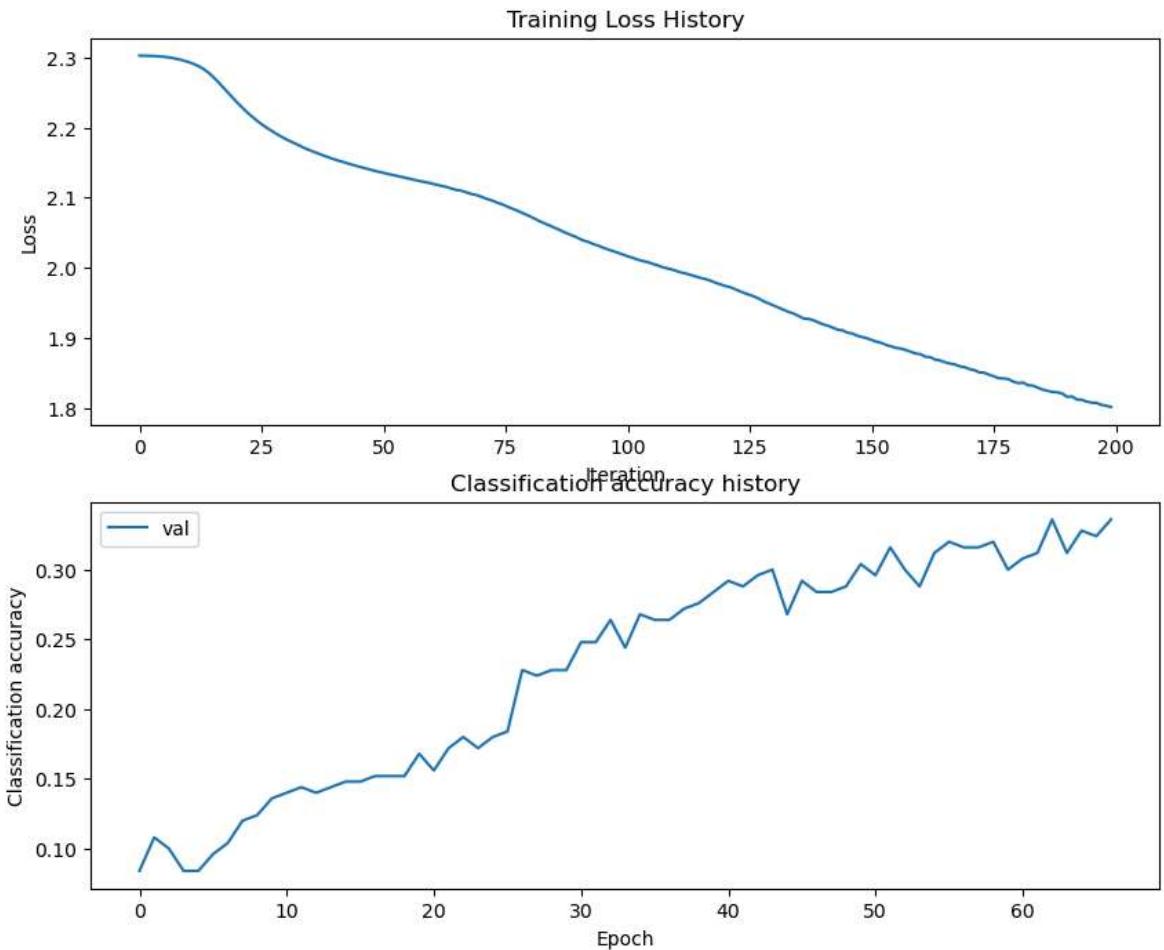
```
In [12]: # Plot the training Loss function and validation accuracies
plt.subplot(2, 1, 1)
```

```

plt.plot(train_error)
plt.title("Training Loss History")
plt.xlabel("Iteration")
plt.ylabel("Loss")

plt.subplot(2, 1, 2)
# plt.plot(stats['train_acc_history'], label='train')
plt.plot(validation_accuracy, label="val")
plt.title("Classification accuracy history")
plt.xlabel("Epoch")
plt.ylabel("Classification accuracy")
plt.legend()
plt.show()

```



```

In [13]: from ece285.utils.vis_utils import visualize_grid

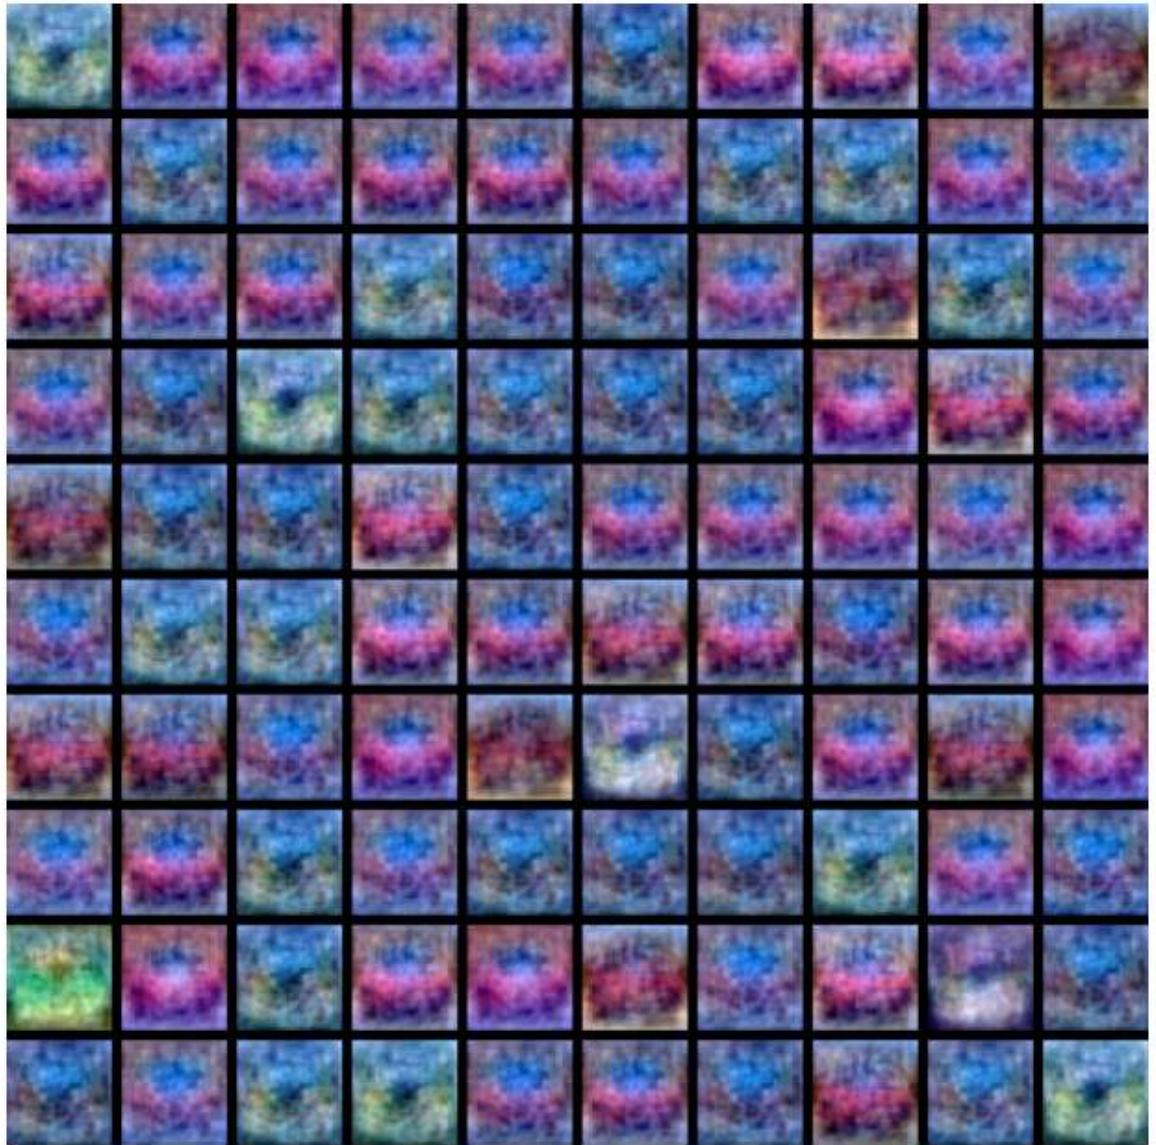
# Credits: http://cs231n.stanford.edu/

# Visualize the weights of the network

def show_net_weights(net):
    W1 = net._modules[0].parameters[0]
    W1 = W1.reshape(3, 32, 32, -1).transpose(3, 1, 2, 0)
    plt.imshow(visualize_grid(W1, padding=3).astype("uint8"))
    plt.gca().axis("off")
    plt.show()

show_net_weights(net)

```



## Tune your hyperparameters (50%)

**What's wrong?** Looking at the visualizations above, we see that the loss is decreasing more or less linearly, which seems to suggest that the learning rate may be too low. Moreover, there is no gap between the training and validation accuracy, suggesting that the model we used has low capacity, and that we should increase its size. On the other hand, with a very large model we would expect to see more overfitting, which would manifest itself as a very large gap between the training and validation accuracy.

**Tuning.** Tuning the hyperparameters and developing intuition for how they affect the final performance is a large part of using Neural Networks, so we want you to get a lot of practice. Below, you should experiment with different values of the various hyperparameters, including hidden layer size, learning rate, numer of training epochs, and regularization strength.

**Approximate results.** You should be aim to achieve a classification accuracy of greater than 40% on the validation set. Our best network gets over 40% on the

validation set.

**Experiment:** Your goal in this exercise is to get as good of a result on cifar10 as you can (40% could serve as a reference), with a fully-connected Neural Network.

## Explain your hyperparameter tuning process below.

Your Answer:

1. Since the loss is decreasing approximate linearly, we can use a higher learning rate to converge faster, therefore I set learning rate to 0.1, a commonly used lr value.
2. There is no gap between the training and validation accuracy, suggesting that the model we used has low capacity. Thus I increase the model size by increasing hidden layer to 300.
3. After tuning above hyperparameters in step 1 and step 2, I found that the training loss converges before 200 epochs. So I set epochs as default 200, and it works well! The validation accuracy goes to 42.8%.

```
In [20]: best_net_hyperparams = None # store the best model into this

#####
# TODO: Tune hyperparameters using the validation set. Store your best trained
# model hyperparams in best_net.
#
# To help debug your network, it may help to use visualizations similar to the
# ones we used above; these visualizations will have significant qualitative
# differences from the ones we saw above for the poorly tuned network.
#
# You are now free to test different combinations of hyperparameters to build
# various models and test them according to the above plots and visualization

input_size = 3072
hidden_size = 300 # Hidden Layer size (Hyper-parameter)
num_classes = 10 # Output

# Initialize the dataset with the dataloader class
dataset = DataLoader(x_train, y_train, x_val, y_val, x_test, y_test)
best_net = init_model()
optim = SGD(best_net, lr = 0.1, weight_decay = 0.01)
loss_func = CrossEntropyLoss()
epoch = 200 # (Hyper-parameter)
batch_size = 200 # (Reduce the batch size if your computer is unable to handle

trainer = Trainer(
    dataset, optim, best_net, loss_func, epoch, batch_size, validate_interval=3
)
train_error, validation_accuracy = trainer.train()

out_train = best_net.predict(x_train)
acc = get_classification_accuracy(out_train, y_train)
print("Training acc: ", acc)
out_val = best_net.predict(x_val)
```

```
acc = get_classification_accuracy(out_val, y_val)
print("Validation acc: ", acc)

# TODO: Show the above plots and visualizations for the default params (already
# done) and the best hyper-params you obtain. You only need to show this for 2
# sets of hyper-params.
# You just need to store values for the hyperparameters in best_net_hyperparams
# as a list in the order
# best_net_hyperparams = [lr, weight_decay, epoch, hidden_size]
#####
best_net_hyperparams = [0.01, 0.01, 200, 100] # default
best_net_hyperparams = [0.1, 0.01, 200, 300]
```

Epoch Average Loss: 2.300071  
Validate Acc: 0.084  
Epoch Average Loss: 2.268512  
Epoch Average Loss: 2.214933  
Epoch Average Loss: 2.179944  
Validate Acc: 0.156  
Epoch Average Loss: 2.167844  
Epoch Average Loss: 2.150929  
Epoch Average Loss: 2.141386  
Validate Acc: 0.212  
Epoch Average Loss: 2.112754  
Epoch Average Loss: 2.092109  
Epoch Average Loss: 2.068048  
Validate Acc: 0.232  
Epoch Average Loss: 2.057241  
Epoch Average Loss: 2.037002  
Epoch Average Loss: 2.045416  
Validate Acc: 0.292  
Epoch Average Loss: 2.041739  
Epoch Average Loss: 2.012062  
Epoch Average Loss: 1.976290  
Validate Acc: 0.264  
Epoch Average Loss: 2.000353  
Epoch Average Loss: 1.983939  
Epoch Average Loss: 1.953881  
Validate Acc: 0.316  
Epoch Average Loss: 1.929595  
Epoch Average Loss: 1.928706  
Epoch Average Loss: 1.909803  
Validate Acc: 0.296  
Epoch Average Loss: 1.904079  
Epoch Average Loss: 1.877448  
Epoch Average Loss: 1.873545  
Validate Acc: 0.276  
Epoch Average Loss: 1.897851  
Epoch Average Loss: 1.834607  
Epoch Average Loss: 1.846808  
Validate Acc: 0.352  
Epoch Average Loss: 1.838307  
Epoch Average Loss: 1.826593  
Epoch Average Loss: 1.827618  
Validate Acc: 0.356  
Epoch Average Loss: 1.814641  
Epoch Average Loss: 1.803634  
Epoch Average Loss: 1.787345  
Validate Acc: 0.328  
Epoch Average Loss: 1.782202  
Epoch Average Loss: 1.803197  
Epoch Average Loss: 1.760699  
Validate Acc: 0.348  
Epoch Average Loss: 1.786442  
Epoch Average Loss: 1.754840  
Epoch Average Loss: 1.777949  
Validate Acc: 0.384  
Epoch Average Loss: 1.792728  
Epoch Average Loss: 1.749026  
Epoch Average Loss: 1.735431  
Validate Acc: 0.360  
Epoch Average Loss: 1.749346  
Epoch Average Loss: 1.754655

Epoch Average Loss: 1.755395  
Validate Acc: 0.332  
Epoch Average Loss: 1.800520  
Epoch Average Loss: 1.733904  
Epoch Average Loss: 1.765536  
Validate Acc: 0.344  
Epoch Average Loss: 1.727875  
Epoch Average Loss: 1.741139  
Epoch Average Loss: 1.724926  
Validate Acc: 0.380  
Epoch Average Loss: 1.737598  
Epoch Average Loss: 1.691864  
Epoch Average Loss: 1.725540  
Validate Acc: 0.380  
Epoch Average Loss: 1.695348  
Epoch Average Loss: 1.707106  
Epoch Average Loss: 1.690490  
Validate Acc: 0.408  
Epoch Average Loss: 1.673739  
Epoch Average Loss: 1.669382  
Epoch Average Loss: 1.723950  
Validate Acc: 0.404  
Epoch Average Loss: 1.706270  
Epoch Average Loss: 1.655192  
Epoch Average Loss: 1.692525  
Validate Acc: 0.340  
Epoch Average Loss: 1.681834  
Epoch Average Loss: 1.655527  
Epoch Average Loss: 1.656362  
Validate Acc: 0.368  
Epoch Average Loss: 1.693376  
Epoch Average Loss: 1.669411  
Epoch Average Loss: 1.665840  
Validate Acc: 0.348  
Epoch Average Loss: 1.688954  
Epoch Average Loss: 1.694815  
Epoch Average Loss: 1.663867  
Validate Acc: 0.372  
Epoch Average Loss: 1.655381  
Epoch Average Loss: 1.639057  
Epoch Average Loss: 1.669988  
Validate Acc: 0.384  
Epoch Average Loss: 1.677361  
Epoch Average Loss: 1.626184  
Epoch Average Loss: 1.623320  
Validate Acc: 0.388  
Epoch Average Loss: 1.643714  
Epoch Average Loss: 1.626226  
Epoch Average Loss: 1.632185  
Validate Acc: 0.368  
Epoch Average Loss: 1.664252  
Epoch Average Loss: 1.617582  
Epoch Average Loss: 1.593336  
Validate Acc: 0.392  
Epoch Average Loss: 1.589056  
Epoch Average Loss: 1.623359  
Epoch Average Loss: 1.685273  
Validate Acc: 0.416  
Epoch Average Loss: 1.645217  
Epoch Average Loss: 1.618991

Epoch Average Loss: 1.606476  
Validate Acc: 0.324  
Epoch Average Loss: 1.614198  
Epoch Average Loss: 1.606709  
Epoch Average Loss: 1.606458  
Validate Acc: 0.388  
Epoch Average Loss: 1.612916  
Epoch Average Loss: 1.625931  
Epoch Average Loss: 1.614406  
Validate Acc: 0.404  
Epoch Average Loss: 1.619226  
Epoch Average Loss: 1.616951  
Epoch Average Loss: 1.599097  
Validate Acc: 0.388  
Epoch Average Loss: 1.593995  
Epoch Average Loss: 1.563755  
Epoch Average Loss: 1.605891  
Validate Acc: 0.400  
Epoch Average Loss: 1.597959  
Epoch Average Loss: 1.617890  
Epoch Average Loss: 1.592449  
Validate Acc: 0.380  
Epoch Average Loss: 1.550429  
Epoch Average Loss: 1.611632  
Epoch Average Loss: 1.577012  
Validate Acc: 0.396  
Epoch Average Loss: 1.555697  
Epoch Average Loss: 1.581042  
Epoch Average Loss: 1.624416  
Validate Acc: 0.384  
Epoch Average Loss: 1.592271  
Epoch Average Loss: 1.600728  
Epoch Average Loss: 1.546571  
Validate Acc: 0.428  
Epoch Average Loss: 1.610162  
Epoch Average Loss: 1.546903  
Epoch Average Loss: 1.578859  
Validate Acc: 0.392  
Epoch Average Loss: 1.608172  
Epoch Average Loss: 1.603163  
Epoch Average Loss: 1.564900  
Validate Acc: 0.372  
Epoch Average Loss: 1.586613  
Epoch Average Loss: 1.534824  
Epoch Average Loss: 1.561804  
Validate Acc: 0.380  
Epoch Average Loss: 1.548545  
Epoch Average Loss: 1.511561  
Epoch Average Loss: 1.601804  
Validate Acc: 0.412  
Epoch Average Loss: 1.592309  
Epoch Average Loss: 1.541224  
Epoch Average Loss: 1.544430  
Validate Acc: 0.432  
Epoch Average Loss: 1.548974  
Epoch Average Loss: 1.541045  
Epoch Average Loss: 1.571556  
Validate Acc: 0.356  
Epoch Average Loss: 1.533488  
Epoch Average Loss: 1.549456

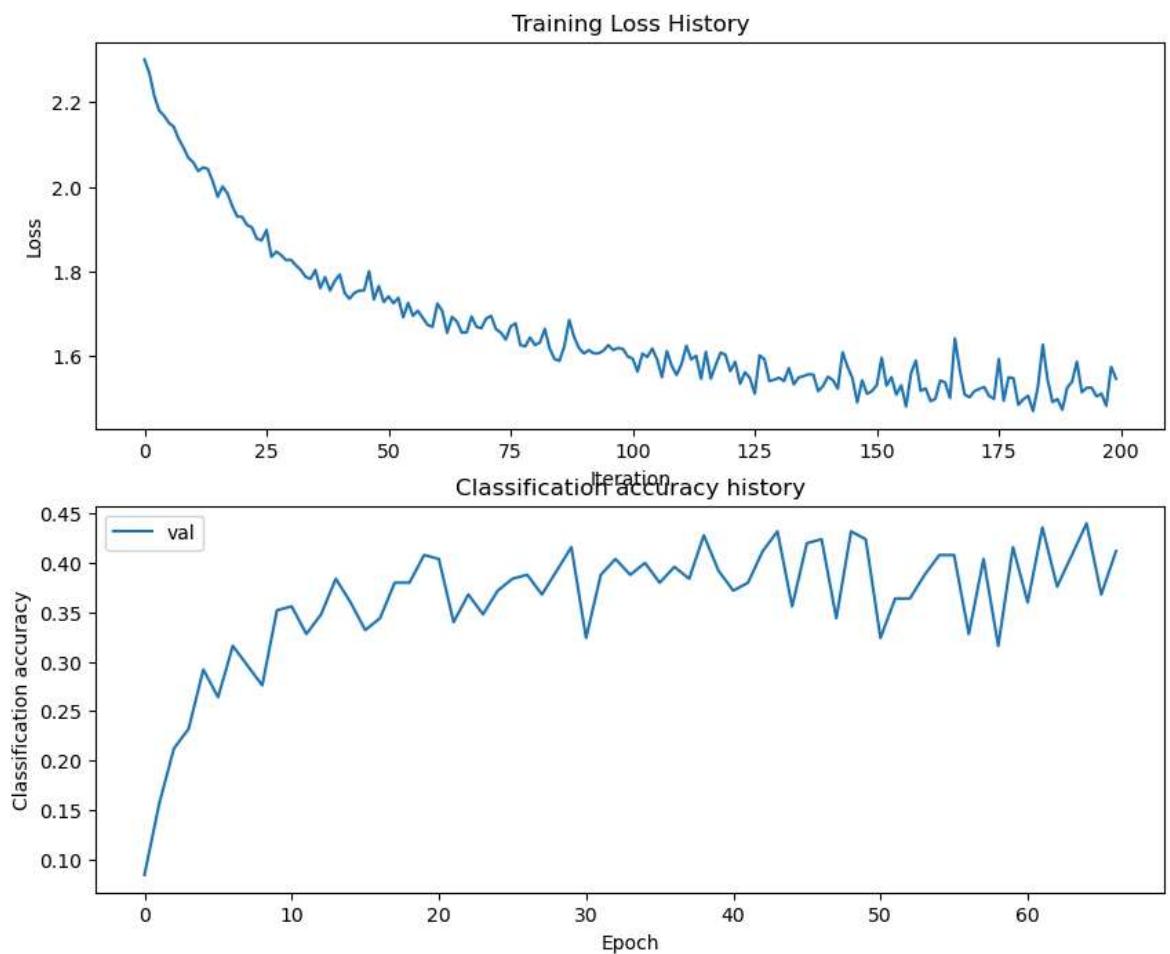
Epoch Average Loss: 1.552967  
Validate Acc: 0.420  
Epoch Average Loss: 1.557000  
Epoch Average Loss: 1.556592  
Epoch Average Loss: 1.517465  
Validate Acc: 0.424  
Epoch Average Loss: 1.529638  
Epoch Average Loss: 1.551317  
Epoch Average Loss: 1.543507  
Validate Acc: 0.344  
Epoch Average Loss: 1.523296  
Epoch Average Loss: 1.609245  
Epoch Average Loss: 1.575333  
Validate Acc: 0.432  
Epoch Average Loss: 1.548461  
Epoch Average Loss: 1.490865  
Epoch Average Loss: 1.542712  
Validate Acc: 0.424  
Epoch Average Loss: 1.511250  
Epoch Average Loss: 1.517506  
Epoch Average Loss: 1.531557  
Validate Acc: 0.324  
Epoch Average Loss: 1.596256  
Epoch Average Loss: 1.530521  
Epoch Average Loss: 1.550064  
Validate Acc: 0.364  
Epoch Average Loss: 1.509095  
Epoch Average Loss: 1.531194  
Epoch Average Loss: 1.480860  
Validate Acc: 0.364  
Epoch Average Loss: 1.558873  
Epoch Average Loss: 1.589553  
Epoch Average Loss: 1.517868  
Validate Acc: 0.388  
Epoch Average Loss: 1.524039  
Epoch Average Loss: 1.493848  
Epoch Average Loss: 1.499255  
Validate Acc: 0.408  
Epoch Average Loss: 1.542604  
Epoch Average Loss: 1.537824  
Epoch Average Loss: 1.501010  
Validate Acc: 0.408  
Epoch Average Loss: 1.641612  
Epoch Average Loss: 1.568103  
Epoch Average Loss: 1.509431  
Validate Acc: 0.328  
Epoch Average Loss: 1.502914  
Epoch Average Loss: 1.517265  
Epoch Average Loss: 1.522744  
Validate Acc: 0.404  
Epoch Average Loss: 1.527083  
Epoch Average Loss: 1.505994  
Epoch Average Loss: 1.499278  
Validate Acc: 0.316  
Epoch Average Loss: 1.593333  
Epoch Average Loss: 1.494689  
Epoch Average Loss: 1.550106  
Validate Acc: 0.416  
Epoch Average Loss: 1.547482  
Epoch Average Loss: 1.485010

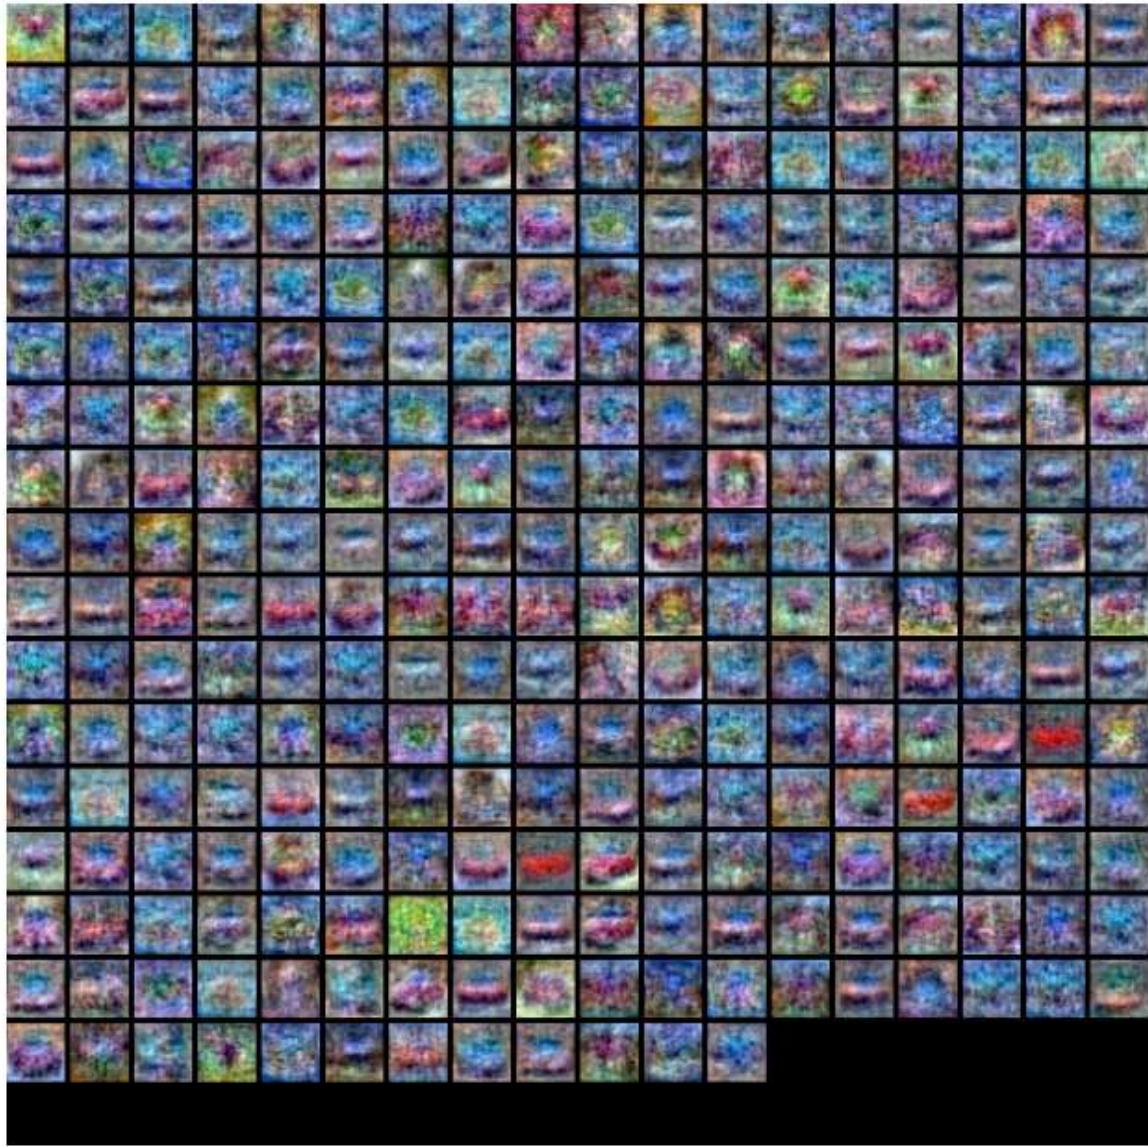
```
Epoch Average Loss: 1.498139
Validate Acc: 0.360
Epoch Average Loss: 1.506377
Epoch Average Loss: 1.470451
Epoch Average Loss: 1.526824
Validate Acc: 0.436
Epoch Average Loss: 1.626705
Epoch Average Loss: 1.543492
Epoch Average Loss: 1.491672
Validate Acc: 0.376
Epoch Average Loss: 1.498718
Epoch Average Loss: 1.473770
Epoch Average Loss: 1.525536
Validate Acc: 0.408
Epoch Average Loss: 1.539913
Epoch Average Loss: 1.586935
Epoch Average Loss: 1.514925
Validate Acc: 0.440
Epoch Average Loss: 1.525352
Epoch Average Loss: 1.525501
Epoch Average Loss: 1.505125
Validate Acc: 0.368
Epoch Average Loss: 1.511678
Epoch Average Loss: 1.482767
Epoch Average Loss: 1.574502
Validate Acc: 0.412
Epoch Average Loss: 1.546934
Training acc: 0.4982
Validation acc: 0.428
```

```
In [21]: # TODO: Plot the training_error and validation_accuracy of the best network (5%)
# Plot the training loss function and validation accuracies
plt.subplot(2, 1, 1)
plt.plot(train_error)
plt.title("Training Loss History")
plt.xlabel("Iteration")
plt.ylabel("Loss")

plt.subplot(2, 1, 2)
# plt.plot(stats['train_acc_history'], label='train')
plt.plot(validation_accuracy, label="val")
plt.title("Classification accuracy history")
plt.xlabel("Epoch")
plt.ylabel("Classification accuracy")
plt.legend()
plt.show()

# TODO: visualize the weights of the best network (5%)
show_net_weights(best_net)
```





## Run on the test set (30%)

When you are done experimenting, you should evaluate your final trained network on the test set; you should get above 35%.

```
In [23]: test_acc = (best_net.predict(x_test) == y_test).mean()  
print("Test accuracy: ", test_acc)
```

Test accuracy: 0.362

### Inline Question (10%)

Now that you have trained a Neural Network classifier, you may find that your testing accuracy is much lower than the training accuracy. In what ways can we decrease this gap? Select all that apply.

1. Train on a larger dataset.
2. Add more hidden units.
3. Increase the regularization strength.
4. None of the above.

### Your Answer:

Both 1 and 3 are correct.

### Your Explanation:

1. Training on a larger dataset is correct because more data helps the model generalize better by exposing it to more variations.
2. Adding more hidden units increases model capacity, which will worsen overfitting (increase the gap between the training and testing accuracy)
3. Increase the regularization strength is correct. Regularization such as L2 weight decay penalizes complex models with large weights. This prevents overfitting.